




DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics: Games Engineering

Hardware-accelerated Raytracing in Dynamic Environments using Vulkan

Christian Mulser





DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics: Games Engineering

Hardware-accelerated Raytracing in Dynamic Environments using Vulkan

Hardwarebeschleunigtes Raytracing in dynamischen Umgebungen mittels Vulkan

Author:	Christian Mulser
Supervisor:	Prof. Dr. Rüdiger Westermann
Advisor:	Prof. Dr. Rüdiger Westermann
Submission Date:	September 19, 2025



I confirm that this bachelor's thesis in informatics: games engineering is my own work and I have documented all sources and material used.

Munich, September 19, 2025

Christian Mulser

Acknowledgments

I thank my supervisor, Rüdiger Westermann, for guiding me along the way and providing useful feedback to my work. I also thank Sebastian Wohner for providing access to the required modern RTX-GPUs and for quickly helping me out with technical difficulties.

Abstract

Contents

Acknowledgments	iii
Abstract	iv
1. Introduction	1
2. Related Work	2
3. Technologies	3
3.1. Operating System	3
3.2. Build Tools	3
3.3. Graphics API	3
3.4. Programming Language	3
3.5. Libraries	4
3.5.1. Volk	4
3.5.2. VMA (Vulkan Memory Allocator)	4
3.5.3. vk-bootstrap	4
3.5.4. glTF-2.0	4
3.5.5. CLI11	5
4. Implementation	6
4.1. Project Setup	6
4.1.1. Source Code Organization	6
4.2. Move Semantics	7
4.3. Shared Pointers	8
4.4. The Builder Pattern	8
4.5. The Vulkan context	9
4.6. The command manager	10
5. Results	12
6. Discussion	13
7. Conclusion	14
A. General Addenda	15
A.1. Detailed Addition	15

B. Figures	16
B.1. Example 1	16
B.2. Example 2	16
List of Figures	17
List of Tables	18
Bibliography	19

1. Introduction

2. Related Work

3. Technologies

3.1. Operating System

Ubuntu 24.04.2 LTS

3.2. Build Tools

- Premake 5
- GNU Make 4.3
- Visual Studio 2022

The project uses Premake 5, which does not build the project directly, but it is a command line utility which reads a scripted definition of a software project and, most commonly, uses it to generate project files for toolsets like Visual Studio, Xcode, or GNU Make [ric23]. It is a simpler alternative to CMake. On Linux we use it to create “Makefiles”, which then ultimately build our project. On windows we create Solution- and Project-files for Visual Studio 2022 and we can build and run the project inside Visual Studio.

3.3. Graphics API

As the graphics API for our project we chose Vulkan 1.2. Vulkan is a modern platform-agnostic API to handle GPU operations. It is the successor to OpenGL, which was discontinued in 2016. Vulkan allows for much more fine-grained control over the GPU, supports multi-threading and for this project most importantly raytracing, through extensions like `VK_KHR_ray_tracing_pipeline` and `VK_KHR_acceleration_structure`. Since these extensions were introduced in Vulkan 1.2, that is the version used in the project.

3.4. Programming Language

At the start of the project Rust was considered as the projects programming language. Rust is a new, memory-safe, low-level programming language and a modern alternative to languages like C and C++. While Rust does have a lot of advantages over C++, we chose C++ in the end. The main reason for this is that the Vulkan API was designed for C/C++. There are a lot of great Vulkan-libraries for Rust, like the popular "Ash" create, which is generated from

"vk.xml". However, way you use the API does differ from the classic C API. This could be problematic, considering that nearly all the few resources relating to raytracing in Vulkan use C/C++. This could have made it difficult to translate the raytracing code into Rust later down the line, so C++ was the safer choice.

3.5. Libraries

3.5.1. Volk

volk is a meta-loader for Vulkan. It allows you to dynamically load entry points required to use Vulkan without linking to vulkan-1.dll or statically linking Vulkan loader. Additionally, volk simplifies the use of Vulkan extensions by automatically loading all associated entry points [<https://github.com/zeux/volk>].

3.5.2. VMA (Vulkan Memory Allocator)

Memory Management in Vulkan is quite difficult. For example, when you create a buffer or image in Vulkan there is no memory associated with it. The programmer needs to manually allocate block of memory with a certain type from a certain memory heap. The properties of these memory heaps depend on the underlying physical device and need to be queried first. Additionally, creating a new memory allocation for each resource may be suboptimal, since the number of allocations can be very limited and many allocations may negatively impact performance. VMA (Vulkan Memory Allocator) simplifies, automates and optimizes memory management in Vulkan.

3.5.3. vk-bootstrap

Initializing Vulkan involves a lot of boilerplate code, which is going to be basically the same in every Vulkan application. VkBootstrap can be used to reduce the amount of boilerplate code in a project. It can simplify the following tasks:

- Instance creation
- Physical device selection
- (Logical) device creation
- Queue acquisition
- Swapchain creation

3.5.4. glTF-2.0

glTF-2.0 is used as the file format for the dynamic 3D scenes that the project renders. It stores its models in a scene graph and supports skeletal animation as well as morph animations. By

default, it stores the scene data in the popular JSON format. Additionally, buffers are stored in binary in BIN files and textures are stored in image files (PNG, JPEG, ...). Since the file format was developed by Khronos Group, which is also behind Vulkan, the data inside the buffers is laid out in a way that makes it easy to use in Vulkan.

To load a glTF file in the project, cgltf is used. cgltf parses the glTF file, which is stored in JSON. It would also have been possible to read the JSON directly, but cgltf takes care of validating the file, loads the required buffers from their BIN-files into memory and creates a C-structure, which is a lot more comfortable to use. Our project does not use this structure directly, but rather used to generate a Scene object.

3.5.5. CLI11

CLI11 is a header-only library written in modern C++, that is designed to parse the command-line options of your application. It is simple, feature-rich and safe. CLI11 allows you to define flags, options, commands and subcommands. You can also add certain restrictions to your parameters, like only allowing positive numbers or that a parameter is an existing path.

4. Implementation

4.1. Project Setup

The projects as well as this thesis' source code are available in the following Github repository:

https://github.com/mulsi/cpp/bachelor_thesis_2025.git

The project is organized in the following directories:

- "assets": This folder contains the projects' assets. It is further divided into scenes, which contains the glTF test scenes, and "shaders", which contains the shaders, that are used in the different pipelines.
- "external": This folder contains all of the "complex" external libraries. "Complex" in this context means, that the library has multiple header or source files or it is in binary form.
- "premake": This folder contains the premake executables for both windows and linux.
- "src": This folder contains all of our source files for the project. It is described in more detail in the next chapter.

Additionally, there are some scripts, that are used to build the project.

4.1.1. Source Code Organization

In this section we will specifically look at how the "src" folder is organised.

Directly in the "src" folder we find "main.cpp", which contains our entry point of the application. The "external" folder once again contains external libraries, but this time the libraries, which are in the form of a single header file. These usually need to be included in a C/C++ file and need to be compiled. That's why it makes sense to keep these libraries with our other source files. The "utils" folder contains utilities, that may be useful all throughout the project. It includes functionalities like debug logging, move semantics and file loading. All the folders with a "vk_" prefix contain code, which abstracts the low-level Vulkan API code. The goal of these abstractions is, that the low-level code should only ever be used in these folders. The rest of the project should only use the abstracted API. Each of the folders focuses on a different area of Vulkan:

- "vk_core": The most important class in this folder is the Context class. It is designed as a singleton, which contains the VkInstance, VkPhysicalDevice, VkDevice, VkQueues and VmaAllocator. After the context gets created it can be accessed anywhere. The

folder also contains the `Handle<T>` class, which ensures that a handle only has one owner, the `CommandBuffer` abstraction and functionality to deal with `VkFormat`.

- "vk_resources": This folder contains Vulkan resource objects like `Buffer` and `Image`, as well as `ImageView` and `SubBuffer`.
- "vk_pipeline": This folder contains Vulkan objects, that are related to the setup of a pipeline. This includes `Pipeline`, `PipelineLayout`, `Shader`, `RenderPass`, `Framebuffer` and `DescriptorPool`.
- "vk_sync": Contains the Vulkan synchronization objects `Fence` and `Semaphore`.
- "scene": This folder contains the `Scene` class, which describes a dynamic scene with a scene graph. All the other classes used by or contained in the scene, like `Mesh`, `Node`, `Camera` and `Animation`, can also be found in this folder.
- "rendering": This folder contains 3 different renderers:
 - `Rasterizer`: Render a scene from the viewpoint of a camera into a framebuffer using traditional rasterisation.
 - `Raytracer`: Render a scene from the viewpoint of a camera into an image using modern hardware-accelerated raytracing.
 - `Skinner`: Calculates the skinned positions of a mesh during an animation.

4.2. Move Semantics

As already mentioned in the `Methods and Technology` chapter, Rust was originally considered as programming language, because it is modern and memory safe. One of Rust's core features is ownership, which means that each value/object can only be owned and managed by only one variable, and when that variable goes out of scope, it alone is responsible for cleaning up the object. When that variable gets assigned to another variable, the new variable becomes the new owner (i.e. the value gets moved). Ownership would be a very nice feature to manage Vulkan objects. Vulkan objects should not be copied on assignment but only moved and when the object is no longer needed (i.e. the owner goes out of scope), it should automatically be destroyed. While C++ does not support ownership directly, it does have so-called move semantics, with which we can emulate ownership. The ownership of Vulkan handles is managed by the `Handle<T>` class. To do this the class declares a default constructor, move constructor and destructor and overrides the move assignment operator. It is also important, that the copy constructor and copy assignment operator are deleted. The reason for that is that copying a Vulkan object is usually very expensive or not even possible and it is almost never necessary. When we want to copy an object, we need to explicitly create a new one first. This avoids unnecessary and complicated accidental creation and copying of new objects.

4.3. Shared Pointers

Since C++11 the language supports so-called smart pointers, which make memory management safer and less complicated. One of these smart pointers is `std::shared_ptr<T>` (or our alias `ptr::Shared<T>`). This type of pointer allocates an object, which can then be shared between multiple owners. The object lives while it has at least one owner, and the last object to go out of scope cleans up the object. This structure can be combined very nicely with our previously defined handle. A Vulkan object can only have one owner, which is a harsh limitation and often causes problems. If we wrap the object in a shared pointer, it can be shared across multiple owners. We also added the class `ToShared<T>`, which a class can inherit from. An instance of that class can then be moved into a shared pointer by calling `to_shared()`.

4.4. The Builder Pattern

As mentioned, before we want the creation of Vulkan objects to be very explicit. The first method to create objects in C++ you think of are constructors, but they prove to be problematic for a couple of reasons. The first problem is that constructors are often called implicitly (e.g. the default constructor gets called when declaring a variable). Default construction only makes sense for very few Vulkan objects (like for example a `VkFence`) and we do not want a Vulkan object to be created when implicitly (or explicitly) calling the default constructor. So, for that reason, the default constructor simply does not create a Vulkan object or, in other words, leaves the handle at `VK_NULL_HANDLE`. Now if a different constructor actually creates an object, the behavior of constructors becomes inconsistent regarding Vulkan object creation. In addition to that, Vulkan objects are almost always created using a `*CreateInfo` struct, which may take a lot of parameters. To solve all of these issues we implemented a `Builder` class for most Vulkan objects. Some exceptions are `Fence` and `ImageView`, which are created with `Fence::create(...)` and `ImageView::create_from(...)` respectively. This makes it very clear when a Vulkan object is created, since it can only be created using a `Builder` or a static method (like `Fence::create(...)`) and especially it can never be created by a constructor. This is also the reason Vulkan objects never have constructor, besides the default constructor.

```
vk::Buffer buffer; // declare a buffer (handle is VK_NULL_HANDLE)

buffer = vk::BufferBuilder()
    .add_usage(VK_BUFFER_USAGE_TRANSFER_DST_BIT)
    .add_usage(VK_BUFFER_USAGE_VERTEX_BUFFER_BIT)
    .memory_usage(VMA_MEMORY_USAGE_GPU_ONLY)
    .size(128)
    .build();
```

4.5. The Vulkan context

Before we can actually do any work in Vulkan, we need to do a lot of setup. This usually includes the following steps:

1. Create a `VkInstance`.
2. Select a `VkPhysicalDevice`, which supports the extensions, features and capabilities required by the project.
3. Create a `VkDevice` from the selected `VkPhysicalDevice`.
4. When creating the `VkDevice`, `VkQueues` from a certain queue family need to be enabled, to be able to submit command buffers later.

These Vulkan objects are accessed very frequently throughout the project. For example, you need to pass the `VkDevice` as a parameter when creating or destroying most Vulkan objects and a command buffer, which executes a sequence of Vulkan commands, needs to be submitted to a certain `VkQueue`. Keeping track of all of these Vulkan objects is quite tedious, which is why we introduced the Vulkan context. It is represented by the `vk::Context` class, which stores these Vulkan objects, allows the programmer to access them anywhere and reduces the effort of setting up Vulkan.

The Vulkan context contains the following objects:

- the `VkInstance`
- the `VkPhysicalDevice` and `VkDevice`
- an optional `GLFWwindow*` from which a `VkSurfaceKHR` is created and also stored (this is necessary for rendering to a window, but it will not be used in our main application, since it is headless)
- the `vk::CommandManager`
- the `VmaAllocator`

The class is designed similar to a singleton, but it deviates in some parts. A singleton usually creates its instance at the start of the application or when it is used for the first time (lazy instantiation). However, `vk::Context` can not be implicitly instantiated this way, because it needs some parameters stored in a `vk::ContextInfo`, when being instantiated. For this reason we need to explicitly create the Vulkan context using `vk::Context::create(...)` with the desired `vk::ContextInfo` before calling `vk::Context::instance()` for the first time or we will get a `nullptr`. While all other objects in our project get cleaned up automatically when going out of scope, the Vulkan context also needs to be destroyed explicitly using `vk::Context::destroy()` when it is no longer needed. If we now want to get the current `VkDevice`, we can call `vk::Context::instance()->get_device()`.

4.6. The command manager

Queue management in Vulkan is a big and complex topic. Vulkan exposes the different queue families on your GPU, which support a subset of the following command types:

- Graphics
- Transfer
- Compute
- Present

When submitting a command buffer to a queue, the queue needs to belong to a queue family, which supports the different types of the commands in the command buffer.

The management of Vulkan queues is completely up to the programmer. You can query the capabilities of the queue families and then enable an arbitrary amount of queues from different queue families. This is especially useful in a multi-threaded application, where you can submit two command buffers to two different queues and they will be executed in parallel.

The `vk::CommandManager` class makes queue management easier. It manages the queue families, as well as their enabled queues and created descriptor pools. Our application is single-threaded in order to avoid too much complexity and because our main objective is to benchmark the performance of dynamic raytracing, which works better in a single-threaded environment. For this reason we also do not need multiple queues or queues designated for a certain type of operations (for example a designated transfer queue). However we do need a queue for every command type, which our command manager ensures. For every command type, it will iterate over all queue families and select the first family, which supports the command type. When creating the device, we enable the first queue of every queue family, which has been selected at least one time for a command type.

The following are some examples of possible queue setups:




Figure 4.1.: Queue setup example with only one queue




Figure 4.2.: Queue setup example with two queues

We also need command pool to be able to allocate command buffers. When creating a command pool, a queue family has to be specified. Because the allocated command buffers can only be submitted to queues of the specified family, we create one command pool for every queue family that is used.

5. Results

6. Discussion

7. Conclusion

A. General Addenda

If there are several additions you want to add, but they do not fit into the thesis itself, they belong here.

A.1. Detailed Addition

Even sections are possible, but usually only used for several elements in, e.g. tables, images, etc.

B. Figures

B.1. Example 1

✓

B.2. Example 2

✗

List of Figures

4.1. Queue setup example with only one queue	10
4.2. Queue setup example with two queues	11

List of Tables

Bibliography

[ric23] ricecookie. *What is Premake?* boh. 2023.